

Problem 1**1(a)**

Given, Loss function,

$$L(y_i, \hat{y}_i) = \log(1 + e^{-y_i \hat{y}_i}) \quad (1)$$

Gradient

$$\begin{aligned} g_i &= \frac{\partial L(y_i, \hat{y}_i)}{\partial \hat{y}_i} \\ &= \frac{\partial}{\partial \hat{y}_i} \log(1 + e^{-y_i \hat{y}_i}) \\ &= \frac{-y_i e^{-y_i \hat{y}_i}}{1 + e^{-y_i \hat{y}_i}} \\ \Rightarrow g_i &= \frac{-y_i e^{-y_i \hat{y}_i}}{1 + e^{-y_i \hat{y}_i}} \end{aligned} \quad (2)$$

1(b)

Let,

$$\gamma^* = \min_{\gamma \in R} \sum_{i=1}^n (-g_i - \gamma h(x_i))^2$$

We solve for γ^* below

$$\begin{aligned} \frac{\partial \sum_{i=1}^n (-g_i - \gamma h(x_i))^2}{\partial \gamma} &= \sum_{i=1}^n 2h(x_i)(g_i + \gamma h(x_i)) = 0 \\ \Rightarrow \sum_{i=1}^n g_i h(x_i) + \gamma \sum_{i=1}^n h(x_i)^2 &= 0 \\ \Rightarrow \gamma^* &= -\frac{\sum_{i=1}^n g_i h(x_i)}{\sum_{i=1}^n h(x_i)^2} \end{aligned}$$

We find Second derivative to check if minimum is achieved

$$\frac{\partial^2 \sum_{i=1}^n (-g_i - \gamma h(x_i))^2}{\partial \gamma^2} = 2 \sum_{i=1}^n h(x_i)^2 > 0$$

Hence, we conclude γ can be computed in a closed form solution. Now we solve for h^*

$$h^* = \min_{h \in R} y^* = \min_{h \in R} -\frac{\sum_{i=1}^n g_i h(x_i)}{\sum_{i=1}^n h(x_i)^2}$$

$$\frac{\partial \frac{\sum_{i=1}^n g_i h(x_i)}{\sum_{i=1}^n h(x_i)^2}}{\partial h} = -\frac{\sum_{i=1}^n g_i h'(x_i)}{\sum_{i=1}^n h(x_i)^2} + \frac{\sum_{i=1}^n g_i h(x_i) \sum_{i=1}^n 2h(x_i) h'(x_i)}{(\sum_{i=1}^n h(x_i)^2)^2} = 0$$

Solving the above equation gives us h^* . Since γ is not at all present in the equation, we can imply that h^* can be derived independent of γ

1(c)

As per Newtons Method, at step $t + 1$, we have,

$$\alpha^{*t+1} = \alpha^{*t} - \frac{f'(\alpha)}{f''(\alpha)} \quad (3)$$

$$f(\alpha) = \sum_{i=1}^n \log(1 + e^{-y_i(\hat{y}_i + \alpha h^*(x_i))})$$

$$\begin{aligned} f'(\alpha) &= \frac{\partial \sum_{i=1}^n \log(1 + e^{-y_i(\hat{y}_i + \alpha h^*(x_i))})}{\partial \alpha} \\ &= \sum_{i=1}^n \frac{-y_i h^*(x_i) e^{-y_i(\hat{y}_i + \alpha h^*(x_i))}}{1 + e^{-y_i(\hat{y}_i + \alpha h^*(x_i))}} \end{aligned}$$

$$\Rightarrow f'(\alpha) = \sum_{i=1}^n \frac{-y_i h^*(x_i)}{e^{y_i(\hat{y}_i + \alpha h^*(x_i))} + 1} \quad (4)$$

$$f''(\alpha) = \frac{\partial f'(\alpha)}{\partial \alpha} = \sum_{i=1}^n \frac{-y_i^2 h^*(x_i)^2 e^{y_i(\hat{y}_i + \alpha h^*(x_i))}}{(e^{y_i(\hat{y}_i + \alpha h^*(x_i))} + 1)^2}$$

$$f''(\alpha) = \sum_{i=1}^n \frac{-y_i^2 h^*(x_i)^2 e^{y_i(\hat{y}_i + \alpha h^*(x_i))}}{(e^{y_i(\hat{y}_i + \alpha h^*(x_i))} + 1)^2} \quad (5)$$

Substituting (4) and (5) in (3), we get, for step $t + 1$,

$$\begin{aligned} \alpha^{*t+1} &= \alpha^* - \frac{f'(\alpha^*)}{f''(\alpha^*)} \\ \Rightarrow \alpha^{*t+1} &= \alpha^* - \sum_{i=1}^n \frac{e^{-y_i(\hat{y}_i + \alpha^* h^*(x_i))} + 1}{y_i h^*(x_i)} \end{aligned} \quad (6)$$

Substituting (6) in the update step, we get,

$$\hat{y}_i = \hat{y}_i + \left[\alpha^* - \sum_{i=1}^n \frac{e^{-y_i(\hat{y}_i + \alpha^* h^*(x_i))} + 1}{y_i h^*(x_i)} \right] h^*(x_i)$$

Problem 2

2(a)

Maximize the flatness means minimize the Norm-2. Hence the Primal optimization formulation is:

$$\text{minimize } \frac{1}{2} \|w\|_2^2 \quad (7)$$

$$y_i - w^T x - b \leq \epsilon \quad (8)$$

$$w^T x + b - y_i \leq \epsilon \quad (9)$$

2(b)

To cope with unfeasible constraints, we can introduce Slack variables ξ_i, ξ_i^* that ensures the deviation above ϵ is taken care of. Hence, the formulation becomes:

$$\text{minimize } \frac{1}{2} \|w\|_2^2 + C \sum_{i=1}^l (\xi_i + \xi_i^*) \quad (10)$$

$$y_i - w^T x - b \leq \epsilon + \xi_i \quad (11)$$

$$w^T x + b - y_i \leq \epsilon + \xi_i^* \quad (12)$$

$$\xi_i, \xi_i^* \geq 0 \quad (13)$$

The Loss function is called ϵ -insensitive loss function and is given by,


$$|\xi|_\epsilon = \begin{cases} 0 & \text{If } |\xi| \leq \epsilon \\ |\xi| - \epsilon & \text{Otherwise} \end{cases} \quad (14)$$

2(c)

Lagrange function for above Primal form is given by :

$$L = \frac{1}{2} \|w\|^2 + C \sum_{i=1}^l (\xi_i + \xi_i^*) - \sum_{i=1}^l (\eta \xi_i + \eta^* \xi_i^*) \\ - \sum_{i=1}^l \alpha_i (\epsilon + \xi_i - y_i + w^T b) - \sum_{i=1}^l \alpha_i^* (\epsilon + \xi_i^* + y_i - w^T b)$$

Here, $\eta_i, \eta_i^*, \alpha_i$ and α_i^* are the Lagrange multipliers. These have to satisfy positivity constraints.

Figure 1: 2.(b) Hinge loss function for $\epsilon = 0.5$

$$\eta_i, \eta_i^*, \alpha_i, \alpha_i^* \geq 0$$

We now take partial derivatives with respect to primal variables.

$$\partial_b L = \sum_{i=1}^l (\alpha_i^* - \alpha_i) = 0 \quad (15)$$

$$\partial_w L = w + \sum_{i=1}^l (\alpha_i^* - \alpha_i) x_i = 0 \quad (16)$$

$$\partial_{\xi_i} L = C - \alpha_i - \eta_i = 0 \quad (17)$$

$$\partial_{\xi_i^*} L = C - \alpha_i^* - \eta_i^* = 0 \quad (18)$$

Substituting (15), (16), (17) and (18) into Lagrange function, we get the dual form:

$$\text{Maximize } -\frac{1}{2} \sum_{i,j=1}^l (\alpha_i - \alpha_i^*)(\alpha_j - \alpha_j^*) \langle x_i, x_j \rangle - \epsilon \sum_{i=1}^l (\alpha_i + \alpha_i^*) + \sum_{i=1}^l y_i (\alpha_i - \alpha_i^*)$$

Subject to

$$\sum_{i=1}^l (\alpha_i - \alpha_i^*) = 0 \text{ and } \alpha_i, \alpha_i^* \in [0, C]$$

Also from (16), $w = \sum_{i=1}^l (\alpha_i - \alpha_i^*) x_i$ and $f(x) = \sum_{i=1}^l (\alpha_i - \alpha_i^*) \langle x_i, x \rangle + b$

2(d)

Given Kernel $k(x, x') := \phi^T(x)\phi(x')$. As the above dual form depends only on the Dot product of x , we can transform dot products to the given Kernel function as below:

$$\text{Maximize} \quad -\frac{1}{2} \sum_{i,j=1}^l (\alpha_i - \alpha_i^*)(\alpha_j - \alpha_j^*)k(x_i, x_j) - \epsilon \sum_{i=1}^l (\alpha_i + \alpha_i^*) + \sum_{i=1}^l y_i(\alpha_i - \alpha_i^*)$$

Subject to

$$\sum_{i=1}^l (\alpha_i - \alpha_i^*) = 0 \text{ and } \alpha_i, \alpha_i^* \in [0, C]$$

Also from (16), $w = \sum_{i=1}^l (\alpha_i - \alpha_i^*)\phi(x_i)$ and $f(x) = \sum_{i=1}^l (\alpha_i - \alpha_i^*)k(x_i, x) + b$

Problem 3**3.1**

After transforming the data, we get Train and Test sets with dimensions 2000 x 45 each.

3.2

Both trainsvm.m and testsvm.m files are implemented. For $C = 0.1$, we observed a test accuracy of 0.9365

3.3

(a)

| C | Average Accuracy(%) | Average Time |
|----------|---------------------|--------------|
| 4^{-6} | 58.65 | 0.8900 |
| 4^{-5} | 90.40 | 0.6067 |
| 4^{-4} | 92.30 | 0.7867 |
| 4^{-3} | 93.70 | 0.7133 |
| 4^{-2} | 94.25 | 0.7567 |
| 4^{-1} | 94.15 | 0.8300 |
| 1 | 93.85 | 0.8600 |
| 4 | 93.65 | 0.9033 |
| 4^2 | 93.55 | 1.0300 |

Increase in C value results in increase in Average accuracy too. But, we observe that the accuracy seems to reach a plateau at some C and starts to decrease. It may be due to

overfitting caused after a particular C , as more weight is put on slack variables.

Increase in C value results in increase of Average Time too. This can be because, as C increases, the model dimension increases, which causes longer compute time.

(b)

We choose C which gives best Average Accuracy. Hence, we choose 4^{-2}

(c)

Test Accuracy = 0.9335

3.4

| C | Average Accuracy(%) | Average Time |
|----------|---------------------|--------------|
| 4^{-6} | 56.50 | 0.2233 |
| 4^{-5} | 91.35 | 0.1867 |
| 4^{-4} | 92.70 | 0.1267 |
| 4^{-3} | 93.60 | 0.0833 |
| 4^{-2} | 94.40 | 0.0633 |
| 4^{-1} | 94.50 | 0.0600 |
| 1 | 94.40 | 0.0600 |
| 4 | 94.40 | 0.1033 |
| 4^2 | 94.55 | 1.1567 |

(a)

As we see from above tables, compared to my SVM, LIBSVM is marginally better overall. But LIBSVM does give better performance after the model hits a plateau at some C .

(b)

LIBSVM is much faster (5-10 times faster on average) compared to my implementation of SVM.

3.5

(a)

Table for Average Accuracy

| C/Degree | 1 | 2 | 3 |
|----------|---------|---------|---------|
| 4^{-3} | 55.7500 | 55.7500 | 55.7500 |
| 4^{-2} | 78.1500 | 78.1500 | 78.1500 |
| 4^{-1} | 92.2500 | 92.2500 | 92.2500 |
| 1 | 94.7500 | 94.7500 | 94.7500 |
| 4 | 96.2000 | 96.2000 | 96.2000 |
| 4^2 | 97.0500 | 97.0500 | 97.0500 |
| 4^3 | 96.9500 | 96.9500 | 96.9500 |
| 4^4 | 96.7500 | 96.7500 | 96.7500 |
| 4^5 | 96.4500 | 96.4500 | 96.4500 |
| 4^6 | 96.4500 | 96.4500 | 96.4500 |
| 4^7 | 96.4500 | 96.4500 | 96.4500 |

Table for Average Time

| C/Degree | 1 | 2 | 3 |
|----------|--------|--------|--------|
| 4^{-3} | 0.2333 | 0.5067 | 0.7633 |
| 4^{-2} | 0.2633 | 0.4933 | 0.7633 |
| 4^{-1} | 0.2067 | 0.3900 | 0.5700 |
| 1 | 0.1133 | 0.2300 | 0.3500 |
| 4 | 0.0800 | 0.1633 | 0.2433 |
| 4^2 | 0.0733 | 0.1433 | 0.2200 |
| 4^3 | 0.0667 | 0.1333 | 0.2000 |
| 4^4 | 0.0667 | 0.1333 | 0.1967 |
| 4^5 | 0.0633 | 0.1333 | 0.1967 |
| 4^6 | 0.0667 | 0.1333 | 0.2033 |
| 4^7 | 0.0667 | 0.1300 | 0.1967 |

(b)

Table for Average Accuracy

| C/Gamma | 4^{-7} | 4^{-6} | 4^{-5} | 4^{-4} | 4^{-3} | 4^{-2} | 4^{-1} |
|----------|----------|----------|----------|----------|----------|----------|----------|
| 4^{-3} | 73.2500 | 73.2500 | 73.2500 | 73.2500 | 73.2500 | 73.2500 | 73.2500 |
| 4^{-2} | 92.3000 | 92.3000 | 92.3000 | 92.3000 | 92.3000 | 92.3000 | 92.3000 |
| 4^{-1} | 93.4000 | 93.4000 | 93.4000 | 93.4000 | 93.4000 | 93.4000 | 93.4000 |
| 1 | 95.1000 | 95.1000 | 95.1000 | 95.1000 | 95.1000 | 95.1000 | 95.1000 |
| 4 | 96.2500 | 96.2500 | 96.2500 | 96.2500 | 96.2500 | 96.2500 | 96.2500 |
| 4^2 | 96.9000 | 96.9000 | 96.9000 | 96.9000 | 96.9000 | 96.9000 | 96.9000 |
| 4^3 | 96.9000 | 96.9000 | 96.9000 | 96.9000 | 96.9000 | 96.9000 | 96.9000 |
| 4^4 | 97.0500 | 97.0500 | 97.0500 | 97.0500 | 97.0500 | 97.0500 | 97.0500 |
| 4^5 | 96.5000 | 96.5000 | 96.5000 | 96.5000 | 96.5000 | 96.5000 | 96.5000 |
| 4^6 | 96.5000 | 96.5000 | 96.5000 | 96.5000 | 96.5000 | 96.5000 | 96.5000 |
| 4^7 | 96.5000 | 96.5000 | 96.5000 | 96.5000 | 96.5000 | 96.5000 | 96.5000 |

Table for Average Time

| C/Gamma | 4^{-7} | 4^{-6} | 4^{-5} | 4^{-4} | 4^{-3} | 4^{-2} | 4^{-1} |
|----------|----------|----------|----------|----------|----------|----------|----------|
| 4^{-3} | 0.2433 | 0.4733 | 0.7067 | 0.9367 | 1.1700 | 1.4067 | 1.6600 |
| 4^{-2} | 0.2000 | 0.4167 | 0.5933 | 0.7733 | 0.9533 | 1.1333 | 1.3167 |
| 4^{-1} | 0.1167 | 0.2333 | 0.3567 | 0.4700 | 0.5967 | 0.7133 | 0.8300 |
| 1 | 0.0900 | 0.1733 | 0.2700 | 0.3633 | 0.4600 | 0.5567 | 0.6533 |
| 4 | 0.0667 | 0.1333 | 0.2067 | 0.2733 | 0.3533 | 0.4200 | 0.4867 |
| 4^2 | 0.0667 | 0.1333 | 0.2000 | 0.2633 | 0.3367 | 0.4000 | 0.4667 |
| 4^3 | 0.0633 | 0.1267 | 0.1867 | 0.2467 | 0.3067 | 0.3600 | 0.4167 |
| 4^4 | 0.0567 | 0.1133 | 0.1700 | 0.2267 | 0.2833 | 0.3433 | 0.3967 |
| 4^5 | 0.0600 | 0.1167 | 0.1733 | 0.2367 | 0.2967 | 0.3533 | 0.4100 |
| 4^6 | 0.0600 | 0.1267 | 0.2067 | 0.2667 | 0.3300 | 0.3967 | 0.4733 |
| 4^7 | 0.0700 | 0.1433 | 0.2100 | 0.2733 | 0.3433 | 0.4133 | 0.4800 |

Comparing outputs from Polynomial and RBF Kernels, we see that 97.05% is the highest accuracy achieved. Polynomial Kernel achieves this accuracy for $C = 4^2$. So we choose:

Kernel: Polynomial

C: 4^2

Degree: 2

Using above parameters in LIBSVM, we get 94.9500% accuracy on Test dataset.